

PEERING THROUGH CENSORSHIP:
TACKLING INTERNET CENSORSHIP USING P2P NETWORKING

A Thesis
submitted to the Faculty of the
Graduate School of Arts and Sciences
of Georgetown University
in partial fulfillment of the requirements for the
degree of
Master of Science
in Computer Science

By

Eliana V. Troper, A.B.

Washington, DC
April 11, 2023

Copyright © 2023 by Eliana V. Troper
All Rights Reserved

**PEERING THROUGH CENSORSHIP:
TACKLING INTERNET CENSORSHIP USING P2P NETWORKING**

Eliana V. Troper, A.B.

Thesis Advisor: Micah Sherr

ABSTRACT

TODO: Optional for master's thesis but these are nice (max 350 words)

INDEX WORDS: [TODO], Theses (academic)

DEDICATION

TODO: Optional

ACKNOWLEDGMENTS

TODO

TABLE OF CONTENTS

ACKNOWLEDGMENTS	v
CHAPTER	
1 Introduction	1
2 Background	3
2.1 Networking concepts	3
2.2 What is censorship circumvention?	4
2.3 Prior techniques	10
2.4 Real world observations	12
2.5 P2P networking	12
3 System design	13
3.1 Theory and goals	13
3.2 Implementation	15
4 System evaluation	16
4.1 Performance	16
4.2 Bootstrapping availability	16
5 Conclusion	17
5.1 Future work	17
APPENDIX	
A Sample appendix	18
A.1 Sample section	18
B Second sample appendix	19
BIBLIOGRAPHY	20

LIST OF FIGURES

2.1	A model of networking, sufficient for exploring censorship circumvention [TODO: Add final arrows, some sample traffic]	5
2.2	Our model of networking, with our adversary model overlaid [TODO: Once networking diagram is finalized, add an adversary]	8

LIST OF TABLES

3.1	Comparison of existing systems.	14
-----	-----------------------------------------	----

CHAPTER 1

INTRODUCTION

People use the internet to communicate and share a variety of content. Some entities choose to censor internet connections. These entities may range from individual network administrators to a full nation-state. In this paper, we explore the state of censorship, existing circumvention systems, and we design and build a system which tackles several of the limitations present in state of the art systems. [LATER: Add xrefs to sections]

We begin by exploring what we mean by censorship circumvention and establishing baseline terminology for this discipline. We then explore the present theory behind approaches to censorship circumvention. We address the ethics of censorship circumvention as a field, and how the field originated and how that affects the current state of the field.

We then propose several ways to classify existing systems to allow for some comparison between them, and we then explore existing and past systems to understand where the field stands, and where limitations occur. We emphasize systems that are either cutting edge or in active use in the real world.

We also provide an overview of the state of peer to peer (P2P) networking, focusing on successful projects operating in the real world. We provide a deeper dive into the Interplanetary File System (IPFS) and the underlying networking library, LibP2P. These programs are used extensively in our proposed system.

Having outlined what is and has succeeded in censorship circumvention in both the real world and theory, we propose a system that addresses the flaws of the current cutting edge systems. We call this system CN. [TODO: fn code name] We discuss the architecture of this system, putting a particular emphasis on the underlying data structure and network topology.

We build this system and build a benchmarking tool and other apps on top of this system. We use this benchmarking tool to show that the system functions with strong performance. We then begin to simulate a censor based upon theoretical and real world adversaries, and then benchmark the system again, showing that even under severely censored circumstances, we are able to maintain [adequate/good/strong] performance. We also explore what regions of the world would have to operate in more censored circumstances at present.

We then conclude and offer some opportunities for future work.

CHAPTER 2

BACKGROUND

The term censorship has taken on a life of it's own in the modern social and political context. We are *not* exploring the general connotation of censorship, rather, we are exploring a specific type of censorship that existing in a networking (particularly internetworking) context. We will explore some key concepts and present a short model of networking.

2.1 NETWORKING CONCEPTS

We are going to consider the internet as it currently exists, rather than historical states of the internet. Some key concepts:

- Internet Protocol (IP) addresses: These are used to represent endpoints of a connection, and are used for routing. Within a network, other identifiers, such as MAC addresses, are used for routing, however, we are focused on routing between networks and will only consider IP based routing as we explore censorship circumvention. We will use IPv4 addresses throughout as sample addresses, though conceptually IPv6 addresses function similarly for the focuses of this paper.

- Domain Name System (DNS): DNS resolution is used to map from human-readable and memorizable addresses that map to IP addresses (e.g. georgetown.edu resolved to 23.185.0.2 at the time of writing). [TODO: Explore the transparency of these, DNSSEC, etc]
- Router: Routers perform routing for networking, as well as other functions that act on network traffic. They may be the home of a NAT, a firewall, or censorship software. The controller of a router may access and manipulate any traffic which passes through it.
- Firewalls: These are software that may restrict traffic based on rules. These rules may be simple or complex. These have positive uses, such as securing a network from attacks.
- Network address translation (NAT): A NAT is used to map network addresses seen on one side of a router to those seen on the other side of a router. These are frequently used to house potentially thousands of devices behind a single internet facing IP address. Most humans use devices located behind a NAT.

2.2 WHAT IS CENSORSHIP CIRCUMVENTION?

When we discuss the subfield of censorship circumvention when discussing networking and security, we are *not* discussing all forms of censorship that a web user might face - we are discussing censorship on-the-wire of active internet traffic, which usually consists of blocking access to specific services. We do not consider a service blocking certain conduct (e.g. a social network restricting hate speech) the censorship that we are circumventing. We aim to allow a user to connect to any service they would like (notwithstanding technical limitations) - not to allow a user to use any service




Figure 2.1: A model of networking, sufficient for exploring censorship circumvention
[TODO: Add final arrows, some sample traffic]

how they would like. Throughout this paper, when we mention censorship, we mean on-the-wire censorship of reaching services.

2.2.1 ETHICS

We recognize that some users run heinous services on the web - most of which are illegal in any existing jurisdiction. We believe that the best approach for ceasing access to these heinous services is best done through a legal system - ie, investigating and identifying the perpetrator(s), turning off the service directly, and prosecuting any perpetrator(s) of said acts. We note that most morally reprehensible acts that are frequently associated with anonymity and censorship circumvention systems involve actions that occur in the real world (e.g. exploitation of children, weapon/drug/human

trafficking all involve actions that occur in the real world either prior to distribution on the internet or triggered by actions on the internet).

The only notable class of crime which can occur without physical, real-world actions is piracy/IP violations. We believe that targetted legal actions (e.g. DMCA takedowns) can mitigate the issue, and that the cost of further mitigation through censorship would lead to large collateral damage of legal content and a chilling effect of legal speech and usage of the web, and as such the possibility of a circumvention system being used for piracy/IP violations is an acceptable risk given the ability to mitigate damage through the legal system.

2.2.2 THREAT MODEL AND ADVERSARY

Censorship circumvention is an abstract concept, and approaching it in a scientific way requires properly defining and scoping the task. We have already defined the task, [TODO: xref] and must now scope the problem. We will scope our problem by creating a **threat model** of the censor we are aiming to thwart, termed **the adversary**.

An adversary has an **area of influence**. This is an area an adversary controls the network of, and can control (either directly or through compulsion) any service providers within. We assume an adversary may have an area of influence consisting of several nations, with at least one nation outside of the area of influence of *any* censor.¹

Our adversary has effectively infinite resources, but is bound by the laws of physics and as such is bounded by the difficulty of computational problems (and the known solutions to these problems, e.g. even if $P=NP$, the solution is not known, and as such

¹This does not mean that this nation does not have the technical capability to censor, but rather does not perform any censorship either directly or via threatening a start of censorship.

cannot be used). These resources may include labor, capital, technical expertise, and total control of the legal system.

We also assume that our adversary has access to the source code of any circumvention system available, and knowledge of the existence of our circumvention system. We discuss why security through obscurity cannot be used in an effective circumvention system in [TODO: xref].

We also assume that our adversary only takes actions **on-the-wire** and at **service providers** within their area of influence - effectively, the adversary isn't interested in any individual in particular and is performing a dragnet operation, rather than an operation targetting a specific entity. We can use a thought experiment to justify this: if an adversary had a person looking over your shoulder or using video surveillance, they could simply look at what you are doing and punish you directly, no need to expend the resources needed to censor the web.

We assume our adversary has some interest in the utility of the internet and certain technologies on the internet. This means two things: an adversary will not shut down the internet in it's entirety² and an adversary has some threshold of unacceptable collateral damage caused by potential censorship techniques.³

Finally, we note that any technique that subverts this adversary will also work when weakening the strength of these assumptions - e.g. if the adversary controls one nation instead of several, we have more area to run a circumvention system from; if

²Internet shutdowns are a trivial way to censor forbidden internet usage - at the expense of shutting down any non-forbidden usage. [TODO: Discuss prior shut downs in another section and xref]

³Similar to a full internet shutdown, if an adversary wanted to block all access to a specific HTTP site, they could block *all* HTTP traffic, at the expense of blocking every HTTP site. This piecemeal web censorship could even be thwarted with a circumvention system that simply proxies HTTP traffic outside of the area of influence within an unblocked protocol and connects to HTTP sites from outside the area of influence. This starts to show that any non-complete shutdown of internet traffic may have gaps for circumvention techniques. We discuss collateral damage further in [TODO: xref].




Figure 2.2: Our model of networking, with our adversary model overlaid [TODO: Once networking diagram is finalized, add an adversary]

the adversary is only willing to take actions at service providers and not on-the-wire, anything that can thwart an adversary that takes both actions could trivially thwart the adversary that only performs actions at one. We also note that a system could detect the type of actions an adversary is taking and maximize performance based upon what subset of possible actions an adversary is taking - e.g. an adversary that is not touching services leaves an easy path for a circumvention system to use - those services themselves.

[TODO: Basic model + adversary]

2.2.3 ORIENTATION OF CENSORSHIP CIRCUMVENTION

The web arose in a relatively haphazard fashion [TODO: discuss a bit more]. The web was not a place where adversaries engaged in censorship (if they paid any mind to the web whatsoever). Instead, as the web became more prominent and more global, individuals realized that they could be tracked, and that they desired anonymity. This led to the development of anonymity software, with Tor being the most prominent example. Censorship circumvention arose as a by-product of goals that were not focused on censorship circumvention. As primitive censorship began (e.g. blocking specific IP addresses of an undesired server), people realized that software that proxies traffic (either in a basic fashion such as a proxy, or a more complex fashion such as Tor) could be used to access services that would otherwise be blocked on a local network.

As this occurred, the web continued to grow in prominence, and the concept of "Web 2.0" arose - regular users could post content to sites such as message boards and social media sites. People could organize quickly and effectively across larger distances and with people they had never met in the past. Governments took notice. These sites began being blocked by these sites. Governments and other entities also saw that users were getting around these blocks by using proxies and anonymity software - and so they gathered the IP addresses of these services and blocked them.

At this point, those who used and built the anonymity software recognized that censorship circumvention would be crucial for maintaining access to these tools, and others realized that censorship circumvention in-and-of itself would be useful. This effectively led to the rise of a subdiscipline that draws heavily from networking, security, and even game theory, but also exists independently from those subdisciplines, with a variety of stakeholders including governments, NGOs, non-profits, independent

programmers, academics, and others interested in developing censorship circumvention techniques. This also leads to a relatively unique feature of this discipline - some techniques were developed in academia, some techniques were developed in the real world, and some were developed in both, and ignoring the results from either thread of research and observation can lead to the development of limited systems.

[TODO: Mention Tor PTs independently of specific methods]

2.2.4 BASIC CENSORSHIP

When internet censorship arose, basic techniques were used. These include:

- Blocking specific IPs: traffic to or from specific
- Blocking DNS resolution
- Legal mechanisms

2.2.5 CURRENT THEORY OF CENSORSHIP CIRCUMVENTION

TODO

[TODO: Mention parrot is dead]

2.3 PRIOR TECHNIQUES

TODO

[TODO: Do mention some things like shadowsocks that are in use but are even further from academics]

2.3.1 [POINT TO POINT OBFUSCATION]

TODO

OBFS AND SCRAMBLESUIT

TODO

[TODO: Others]

2.3.2 [P2P AND P2P-LIKE]

TODO

SNOWFLAKE

TODO

[TODO: Iran case study]

MASSBROWSER

TODO

2.3.3 HIGH-RELIABILITY CHANNELS

TODO

DOMAIN FRONTING

TODO

RAVEN

TODO

2.3.4 DECOY ROUTING

TODO

2.4 REAL WORLD OBSERVATIONS

TODO

2.4.1 REAL WORLD VS. ACADEMIC THEORY

TODO

2.5 P2P NETWORKING

TODO

2.5.1 IPFS

TODO

CHAPTER 3

SYSTEM DESIGN

We saw that several circumvention systems work given the constraints they were designed for, however, these constraints are either too loose and have known real-world failures under our adversary model, the systems are too difficult to deploy, or they have hit scalability issues.

We propose a system with a clear theoretical backing and achievable goals based on both theory and real-world observations.

3.1 THEORY AND GOALS

First, we want to define system goal and highlight some non-goals that are close to censorship circumvention, but are not directly used for them.

We aim to thwart the adversary we outline in [xref]. We will rely heavily on the concept of collateral damage. [TODO: Discuss collateral damage further]

We are *not* aiming to provide anonymity.¹ We feel the need to emphasize this, as most current systems that are functional in the real world are used as pluggable transports by Tor.²

¹Perhaps our system could be used to access an anonymity system, but we do not want to guarantee anonymity or guarantee that the system will not leak information that may break anonymity, though we are not deliberately trying to break anonymity, we are also not evaluating our system for it at this point.

²Our system could be used as a PT, as we do have a golang implementation, however, evaluating this system for usage as a PT is reserved for future work.

Table 3.1: Comparison of existing systems.

[TODO: Make table take whole page sideways, make in latex when finalized]

		Approach			Bootstrapping			Limitations				Active usage	Source	Notes [TODO: Make footnote]
		Look-like-nothing	Maximize collateral damage	Refraction Networking	Domain fronting	Email	Out of band	Bridge distribution problem	Scalability	Deployment Difficulty	Low bandwidth			
System	obfs4	X (Custom)			X	X	X	X				PT (Tor)	[TODO]	
	Snowflake		P2P-like (WebRTC)		X				X			PT (Tor)	[TODO]	
	MassBrowser [Refraction Networking]		P2P [TODO]					[TODO]	[TODO]	[TODO]	[TODO]		[TODO]	
				X [TODO]	[TODO]	[TODO]	[TODO]		X	X			[TODO]	
	Meek		Domain fronting (HTTPS)		X				X		*	PT (Tor)	[TODO]	Poor scalability causes low bandwidth
	Raven [TODO: More]		Realistic behavior (Email)			X					X		[TODO]	

3.1.1 COMPARING EXISTING SYSTEMS

Earlier, [TODO: xref] we discussed a variety of censorship circumvention systems. Observing how and why each of these work, and the drawbacks associated with each of them will allow us to formulate goals we wish to achieve and working techniques we may use. We present a short summary of previous systems in [TODO: xref table], comparing and contrasting their successful techniques and transports, and their individual drawbacks.

We see that existing systems are limited by at least one of the following:

- The bridge distribution problem [TODO: Xref]
- Scalability
- Deployment difficulty
- Low bandwidth

Another key observation is that *collateral damage* is the bedrock of the theory behind a circumvention system - as such, any system must be based around maximizing this - the ideal system will make the only feasible way to shut down the system effectively shutting down a large portion of internet traffic.

[TODO: Discuss false positives + the overhead of blocking traffic]

[TODO: More]

We also note that existing systems require a user to download and use specialized software, and that systems usually require using anonymity software such as Tor to reap the benefits of censorship circumvention. While anonymity is a useful [thing] to have in some circumstances on the web, getting the drawbacks of common anonymity software [TODO: outline] is not ideal.

3.2 IMPLEMENTATION

TODO

CHAPTER 4

SYSTEM EVALUATION

TODO

4.1 PERFORMANCE

TODO

4.2 BOOTSTRAPPING AVAILABILITY

TODO

CHAPTER 5

CONCLUSION

TODO

5.1 FUTURE WORK

TODO

APPENDIX A

SAMPLE APPENDIX

SAMPLE

A.1 SAMPLE SECTION

SAMPLE

APPENDIX B

SECOND SAMPLE APPENDIX

SAMPLE

BIBLIOGRAPHY